

CS344

Assignment-1

Aditya Patidar 200101009

Aman Soni 200101012

Kartik Malasiya 200101053

Part 1- *Kernel Threads*

The following three system calls have been created in this part:

- 1) `thread_create()`
- 2) `thread_join()`
- 3) `thread_exit()`

thread_create() -- This call create a new kernel thread which shares the address space with the calling space.

thread_join() -- This call waits for a child thread that shares the address space with the calling process. It will return the PID of the child or -1 if there is none.

thread_exit() -- This call allows a thread to terminate.

The following xv6 OS files need to be changed in order to implement the above functions.

- ➔ `proc.c`
- ➔ `syscall.c`
- ➔ `syscall.h`
- ➔ `defs.h`
- ➔ `sysproc.c`
- ➔ `user.h`
- ➔ `usys.S`

Now we will change the above files individually.

1. *proc.c*-

The functions mentioned above(`thread_create()`, `thread_join()`, `thread_exit()`) are added in this file.

```

int thread_create(void(*fcn)(void*),void* arg,void* stack){
    int i,pid;
    struct proc *np;
    struct proc *curproc = myproc();

    if((np= allocproc()) == 0){
        return -1;
    }
    np->pgdir = curproc->pgdir;
    np->sz = curproc->sz;
    np->parent = curproc;
    *np->tf = *curproc->tf;

    np->tf->eax = 0;
    np->tf->eip = (uint)fcn;
    np->tf->esp = (uint)stack+4069;

    np->tf->esp -= 4;
    *(uint*)(np->tf->esp) = (uint)(arg);

    np->tf->esp -= 4;
    *(uint*)(np->tf->esp) = (uint)0xFFFFFFFF;
    np->tf->eax = 0;

    for(i = 0; i < NOFILE; i++)
        if(curproc->ofile[i])
            np->ofile[i] = filedup(curproc->ofile[i]);
    np->cwd = idup(curproc->cwd);

    safestrcpy(np->name, curproc->name, sizeof(curproc->name));

    pid = np->pid;

    acquire(&ptable.lock);

    np->state = RUNNABLE;

    release(&ptable.lock);

    return pid;
}

```

```

int
thread_join(void)
{
    struct proc *p;
    int havekids, pid;
    struct proc *curproc = myproc();

    acquire(&ptable.lock);
    for(;;){
        havekids = 0;
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->parent != curproc)
                continue;
            havekids = 1;
            if(p->state == ZOMBIE){
                pid = p->pid;
                kfree(p->kstack);
                p->kstack = 0;
                p->pid = 0;
                p->parent = 0;
                p->name[0] = 0;
                p->killed = 0;
                p->state = UNUSED;
                release(&ptable.lock);
                return pid;
            }
        }

        if(!havekids || curproc->killed){
            release(&ptable.lock);
            return -1;
        }

        sleep(curproc, &ptable.lock);
    }
}

```

```

void
thread_exit(void)
{
    struct proc *curproc = myproc();
    struct proc *p;
    int fd;

    if(curproc == initproc)
        panic("init exiting");

    for(fd = 0; fd < NOFILE; fd++){
        if(curproc->ofile[fd]){
            fileclose(curproc->ofile[fd]);
            curproc->ofile[fd] = 0;
        }
    }

    begin_op();
    iput(curproc->cwd);
    end_op();
    curproc->cwd = 0;

    acquire(&ptable.lock);

    wakeup1(curproc->parent);

    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->parent == curproc){
            p->parent = initproc;
            if(p->state == ZOMBIE)
                wakeup1(initproc);
        }
    }
}

```

2. *syscall.c-*

The following lines of the code as shown in the screenshot below have been added to the file to make the system understand that these are also the system calls.

```

extern int sys_thread_create(void);
extern int sys_thread_join(void);
extern int sys_thread_exit(void);

```

3. *syscall.h*

The file syscall.h has been edited to add three new system call macros to the list of already existing system call macros with their respective IDs. The screenshot of the same is given below.

```
2 #define SYS_fork      1
3 #define SYS_exit      2
4 #define SYS_wait      3
5 #define SYS_pipe      4
6 #define SYS_read       5
7 #define SYS_kill       6
8 #define SYS_exec       7
9 #define SYS_fstat      8
10 #define SYS_chdir     9
11 #define SYS_dup      10
12 #define SYS_getpid    11
13 #define SYS_sbrk      12
14 #define SYS_sleep     13
15 #define SYS_uptime    14
16 #define SYS_open      15
17 #define SYS_write     16
18 #define SYS_mknod     17
19 #define SYS_unlink    18
20 #define SYS_link      19
21 #define SYS_mkdir     20
22 #define SYS_close     21
23 #define SYS_thread_create 22
24 #define SYS_thread_join 23
25 #define SYS_thread_exit 24
```

4. *defs.h*

As we know file defs.h contains the prototype for all the system calls and we have created three new system calls so we will have to add prototype for these system calls in defs.h. The screenshot for the same is attached below.

```

105 // proc.c
106 int      cpuid(void);
107 void     exit(void);
108 int      fork(void);
109 int      growproc(int);
110 int      kill(int);
111 struct cpu* mycpu(void);
112 struct proc* myproc();
113 void     pinit(void);
114 void     procdump(void);
115 void     scheduler(void) __attribute__((noreturn));
116 void     sched(void);
117 void     setproc(struct proc*);
118 void     sleep(void*, struct spinlock*);
119 void     userinit(void);
120 int      wait(void);
121 void     wakeup(void*);
122 void     yield(void);
123 //user defined below
124 int      thread_create(void*)(void*),void*,void*);
125 int      thread_join(void);
126 void     thread_exit(void);

```

5. *sysproc.c*

These lines of code are the function corresponding to the system call made by us that will call the actual function of `thread_exit()`, `thread_join()` and `thread_exit()` when the system calls one of these functions.

```

93 int
94 sys_thread_create(void){
95     int fcn;
96     char* arg;
97     char* stack;
98     if(argint(0,&fcn)<0)
99         return -1;
100     if(argstr(1,&arg) <0)
101         return -1;
102     if(argstr(2,&stack)<0)
103         return -1;
104     return thread_create((void*)(void*))fcn,arg,stack);
105 }
106
107 int sys_thread_exit(void){
108     thread_exit();
109     return 0;
110 }
111
112 int sys_thread_join(void){
113     return thread_join();
114 }

```

6. *user.h*

This file also contains prototypes for the system calls. Given below is the screenshot of the added lines.

```
26 int thread_create(void (*)(void*), void*, void*);  
27 int thread_join(void);  
28 void thread_exit(void);
```

7. *usys.S*

This file contains code corresponding to binding the function with the system.

```
32 SYSCALL(thread_create)  
33 SYSCALL(thread_join)  
34 SYSCALL(thread_exit)
```

Part 2: Synchronization

By creating spinlocks and mutexes, we must focus on the synchronisation component in this section.

Two balance structures, b1 and b2, with initial values of 3200 and 2800 each, have been built. And by using the do work function in both of them, we created two

threads to update the shared balance variable. However, if they don't execute with synchronisation, the value of the shared balance doesn't come out as expected (6000), so in order to fix this problem, we had to create spinlocks. However, the problem with spinlocks is that they can't operate properly on a single processor system or when the system is under a lot of load because all the spinlocks will spin endlessly. So, to solve this problem we are going to use mutexes in place of spin locks. The results of five time calls made without synchronisation.

To solve this we will use mutex locks.

Without using mutex locks-

```
Booting from Hard Disk..xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap sta8
init: starting sh
$ thread
```

```
Done s:2F9C
Done s:2F78
Threads finished: (4):5, (5):4, shared balance:3200
$ thread
```

```
Done s:2F9C
Done s:2F78
Threads finished: (7):8, (8):7, shared balance:3200
```

with using mutex locks-

```
$ thread
Starting do_work:s:b2
k:s:b1
Done s:2F78
Done s:2F9C
Threads finished: (5):5, (6):6,shared balance:6000
$ thread
Starting do_work:s:b1
Starting do_work:s:b2
Done s:2F78
Done s:2F9C
Threads finished: (8):8, (9):9,shared balance:6000
```

Thread.c file which is created in xv6 is given below-

Thread.c

```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4 #include "x86.h"
5
6 struct balance{
7     char name[32];
8     int amount;
9 };
10
11 volatile int total_balance = 0;
12
13 struct thread_spinlock lock;
14 struct mutex_lock m_lock;
15
16
17 volatile unsigned int delay(unsigned int d){
18     unsigned int i;
19     for(i=0;i<d;i++){
20         __asm volatile("nop" :::);
21     }
22     return i;
23 }
24
25 struct thread_spinlock{
26     uint locked;
27
28     char* name;
29 };
30
31 void thread_initloc(struct thread_spinlock *lk, char *name){
32     lk->name = name;
33     lk->locked = 0;
34 }
35
36 void thread_spin_lock(struct thread_spinlock* lk){
37     while(xchg(&lk->locked,1)!=0);
38     __sync_synchronize();
39 }
40
41 void thread_spin_unlock(struct thread_spinlock* lk){
42     __sync_synchronize();
43     asm volatile("movl $0, %0" : "+m" (lk->locked) :);
44 }
45
46 struct mutex_lock{
47     uint locked;
```

```

48 };
49
50 void mutex_initlock(struct mutex_lock* lk){
51     lk->locked = 0;
52 }
53
54 void mutex_lock(struct mutex_lock* lk){
55     while(xchg(&lk->locked,1)!=0)
56         sleep(1);
57     __sync_synchronize();
58 }
59
60 void mutex_unlock(struct mutex_lock* lk){
61     __sync_synchronize();
62
63     asm volatile("movl $0, %0" : "+m" (lk->locked) : );
64 }
65
66 void do_work(void*arg){
67     int i;
68     int old;
69
70     struct balance *b = (struct balance*)arg;
71     printf(1,"Starting do_work:s:%s\n",b->name);
72
73     for(i=0;i<b->amount;i++){
74         //thread_spin_lock(&lock);
75         mutex_lock(&m_lock);
76         old = total_balance;
77         delay(100000);
78         total_balance = old + 1;
79         mutex_unlock(&m_lock);
80         //thread_spin_unlock(&lock);
81     }
82     printf(1,"Done s:%x\n",b->name);
83
84     thread_exit();
85     return;
86 }
87
88 int main(int argc,char* argv[]){
89     struct balance b1 = {"b1",3200};
90     struct balance b2 = {"b2",2800};
91
92     void*s1,*s2;
93     int t1,t2,r1,r2;
94
95     //thread_init(&m_lock);
96     mutex_initlock(&m_lock);
97
98     s1=malloc(4096);
99     s2=malloc(4096);
100
101     t1=thread_create(do_work,(void*)&b1,s1);
102     t2=thread_create(do_work,(void*)&b2,s2);
103
104     r1 = thread_join();
105     r2 = thread_join();
106
107     printf(1,"Threads finished: (%d):%d, (%d):%d,shared balance:%d\n",t1,r1,t2,r2,total_balance);
108
109     exit();
110 }
111

```

The call relating to the thread.c has been added to include this in the list of xv6 commands in this makefile.

This line of code has been added to UPROGS.

```
168 UPROGS=\
169     _cat\
170     _echo\
171     _forktest\
172     _grep\
173     _init\
174     _kill\
175     _ln\
176     _ls\
177     _mkdir\
178     _rm\
179     _sh\
180     _stressfs\
181     _wc\
182     _abc\
183     _zombie\
184     _thread\
```

```
251 EXTRA=\
252     mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
253     ln.c ls.c mkdir.c rm.c stressfs.c wc.c zombie.c\
254     printf.c umalloc.c\
255     README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
256     .gdbinit.tmpl gdbutil\
```

How to run the test case in steps

Enter the xv6-public directory first, then issue the make clean command. make command, followed by make qemu-nox command. After that, execute the test case thread.

Remark: When we use mutexes, we achieve the intended results, but when we don't, we get the incorrect results.